

Blood Donations

Bridget Bertoni

April 12, 2018

```
# read in data
train_data=read.csv(file="/Users/bbertoni/Desktop/github/Blood_Donations/training_data.csv",header=T)
test_data=read.csv(file="/Users/bbertoni/Desktop/github/Blood_Donations/test_data.csv",header=T)

names(train_data)=c("ID","Last_donation","Num_donations","Volume","First_donation","Made_donation")
names(test_data)=c("ID","Last_donation","Num_donations","Volume","First_donation")

head(train_data)

##      ID Last_donation Num_donations Volume First_donation Made_donation
## 1  619              2             50  12500             98             1
## 2  664              0             13   3250             28             1
## 3  441              1             16   4000             35             1
## 4  160              2             20   5000             45             1
## 5  358              1             24   6000             77             0
## 6  335              4              4   1000              4             0

head(test_data)

##      ID Last_donation Num_donations Volume First_donation
## 1  659              2             12   3000             52
## 2  276             21              7   1750             38
## 3  263              4              1    250              4
## 4  303             11             11   2750             38
## 5   83              4             12   3000             34
## 6  500              3             21   5250             42

#train_data$Made_donation=as.factor(train_data$Made_donation)
#test_data$Made_donation=as.factor(test_data$Made_donation)

# split training data into a training set and a validation set
set.seed(333)
train=sample(1:nrow(train_data),0.7*nrow(train_data),replace=F)
val_data = train_data[-train,]
train_data = train_data[train,]

# check for missing or strange values
sum(is.na(train_data))

## [1] 0

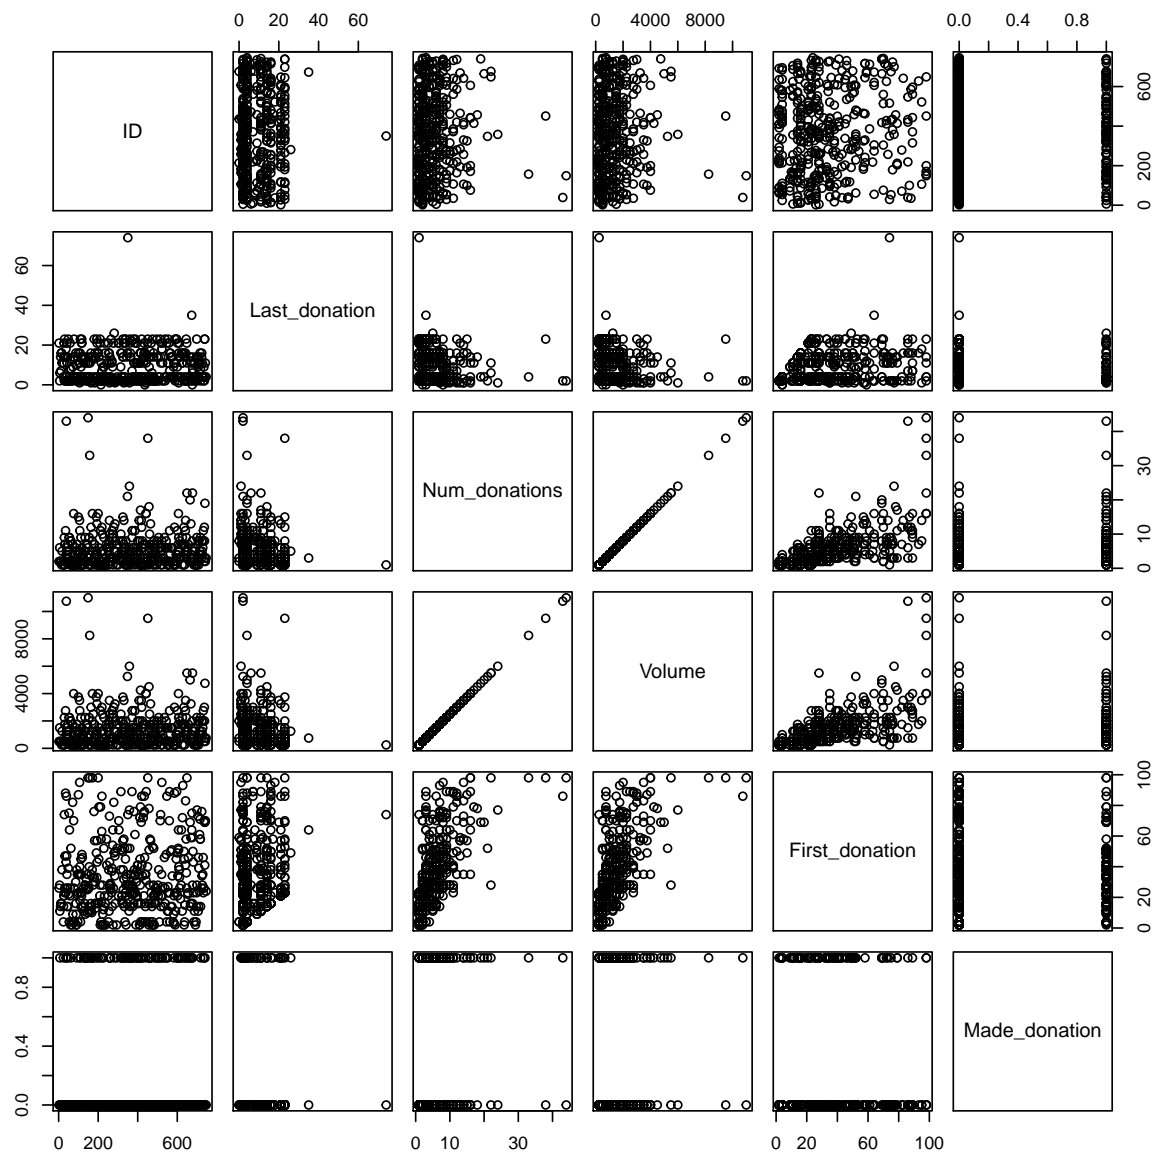
sum(is.na(val_data))# no missing values

## [1] 0





sum(is.na(test_data)) # missing values

## [1] 0



plot(train_data)
```



```
par(mfrow=c(2,2))
plot(as.factor(train_data$Made_donation),train_data$Last_donation)
plot(as.factor(train_data$Made_donation),train_data$Num_donations)
plot(as.factor(train_data$Made_donation),train_data$First_donation)
plot(as.factor(train_data$Made_donation),train_data$Last_donation*train_data$Last_donation)
```



```
plot(as.factor(train_data$Made_donation),train_data$Num_donation*train_data$Num_donation)
plot(as.factor(train_data$Made_donation),train_data$First_donation*train_data$First_donation)
par(mfrow=c(1,1))
```



```
cor(train_data) # correlation between volume and the number of donations is 1
```

```
##              ID Last_donation Num_donations      Volume
## ID          1.00000000    0.02180929    0.02167073    0.02167073
## Last_donation 0.02180929    1.00000000   -0.16150193  -0.16150193
## Num_donations 0.02167073   -0.16150193    1.00000000    1.00000000
## Volume        0.02167073   -0.16150193    1.00000000    1.00000000
## First_donation 0.10391260    0.17791940    0.64740429    0.64740429
## Made_donation 0.04101195   -0.24721682    0.20707525    0.20707525
##
##      First_donation Made_donation
## ID          0.103912600   0.041011949
## Last_donation 0.177919401  -0.247216821
## Num_donations 0.647404293   0.207075245
## Volume        0.647404293   0.207075245
```

```

## First_donation      1.000000000    0.008180282
## Made_donation       0.008180282    1.000000000

# fit basic logistic regression
glm.fit=glm(Made_donation~Last_donation+Num_donations+First_donation,data=train_data,
            family=binomial)
summary(glm.fit)

##
## Call:
## glm(formula = Made_donation ~ Last_donation + Num_donations +
##      First_donation, family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3459  -0.7959  -0.5454  -0.3334   2.4440
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -0.715501   0.240869  -2.970  0.00297 **
## Last_donation -0.082632   0.021242  -3.890  0.00010 ***
## Num_donations  0.109057   0.033638   3.242  0.00119 **
## First_donation -0.012572   0.007533  -1.669  0.09513 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 440.16  on 402  degrees of freedom
## Residual deviance: 396.46  on 399  degrees of freedom
## AIC: 404.46
##
## Number of Fisher Scoring iterations: 5

#plot(glm.fit)
glm.probs=predict(glm.fit,val_data,type="response")
glm.pred=rep(0,nrow(val_data))
glm.pred[glm.probs>0.5]=1
table(glm.pred,val_data$Made_donation) # confusion matrix

##
## glm.pred    0    1
##      0 129  38
##      1   1   5

mean(glm.pred==val_data$Made_donation)

## [1] 0.7745665

# log loss
-(1/nrow(val_data))*( sum(val_data$Made_donation*log(glm.probs)) +
  sum((1-val_data$Made_donation)*log(1-glm.probs)) )

## [1] 0.4671785

# calculate predictions using all of the training data
final_train_data=rbind(train_data,val_data)

```

```

glm.fit=glm(Made_donation~Last_donation+Num_donations+First_donation,data=train_data,
            family=binomial)
glm.probs=predict(glm.fit,test_data,type="response")
out=data.frame(test_data$ID,glm.probs)
names(out)=c("", "Made Donation in March 2007")
write.csv(out,file="logreg_2018_04_16.csv",row.names=FALSE)

# fit basic logistic regression, drop first donation
glm.fit=glm(Made_donation~Last_donation+Num_donations,data=train_data,
            family=binomial)
summary(glm.fit)

```

```

##
## Call:
## glm(formula = Made_donation ~ Last_donation + Num_donations,
##      family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.0754  -0.8127  -0.5449  -0.3272   2.4324
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -0.82025    0.23322  -3.517 0.000436 ***
## Last_donation -0.09346    0.02029  -4.606 4.1e-06 ***
## Num_donations  0.06903    0.02210   3.124 0.001785 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 440.16  on 402  degrees of freedom
## Residual deviance: 399.42  on 400  degrees of freedom
## AIC: 405.42
##
## Number of Fisher Scoring iterations: 5

```

```

#plot(glm.fit)
glm.probs=predict(glm.fit,val_data,type="response")
glm.pred=rep(0,nrow(val_data))
glm.pred[glm.probs>0.5]=1
table(glm.pred,val_data$Made_donation) # confusion matrix

```

```

##
## glm.pred    0    1
##           0 128  39
##           1   2   4

```

```

mean(glm.pred==val_data$Made_donation)

```

```

## [1] 0.7630058

```

```

# log loss
-(1/nrow(val_data))*( sum(val_data$Made_donation*log(glm.probs)) +
  sum((1-val_data$Made_donation)*log(1-glm.probs)) )

```

```
## [1] 0.4820182
```

```
# calculate predictions using all of the training data
```

```
final_train_data=rbind(train_data,val_data)
```

```
glm.fit=glm(Made_donation~Last_donation+Num_donations,data=train_data,  
            family=binomial)
```

```
glm.probs=predict(glm.fit,test_data,type="response")
```

```
out=data.frame(test_data$ID,glm.probs)
```

```
names(out)=c("", "Made Donation in March 2007")
```

```
write.csv(out,file="logreg_nofirstdon_2018_04_16.csv",row.names=FALSE)
```

```
# fit logistic regression with interaction terms
```

```
glm.fit=glm(Made_donation~Last_donation+Num_donations+First_donation+Last_donation:Num_donations,data=t.  
            family=binomial)
```

```
summary(glm.fit)
```

```
##
```

```
## Call:
```

```
## glm(formula = Made_donation ~ Last_donation + Num_donations +
```

```
## First_donation + Last_donation:Num_donations, family = binomial,
```

```
## data = train_data)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -2.6845  -0.7724  -0.5472  -0.3417   2.4294
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)    -0.881040   0.290702  -3.031  0.00244 **  
## Last_donation    -0.062340   0.028078  -2.220  0.02640 *  
## Num_donations     0.140875   0.046997   2.998  0.00272 **  
## First_donation   -0.013458   0.007744  -1.738  0.08221 .  
## Last_donation:Num_donations -0.003388   0.003301  -1.026  0.30473
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
```

```
## Null deviance: 440.16 on 402 degrees of freedom
```

```
## Residual deviance: 395.29 on 398 degrees of freedom
```

```
## AIC: 405.29
```

```
##
```

```
## Number of Fisher Scoring iterations: 5
```

```
#plot(glm.fit)
```

```
glm.probs=predict(glm.fit,val_data,type="response")
```

```
glm.pred=rep(0,nrow(val_data))
```

```
glm.pred[glm.probs>0.5]=1
```

```
table(glm.pred,val_data$Made_donation) # confusion matrix
```

```
##
```

```
## glm.pred    0    1
```

```
##           0 126 36
```

```
##           1   4  7
```

```

mean(glm.pred==val_data$Made_donation)

## [1] 0.7687861

# log loss
-(1/nrow(val_data))*( sum(val_data$Made_donation*log(glm.probs)) +
  sum((1-val_data$Made_donation)*log(1-glm.probs)) )

## [1] 0.4631748

# calculate predictions using all of the training data
final_train_data=rbind(train_data,val_data)
glm.fit=glm(Made_donation~Last_donation+Num_donations+First_donation+Last_donation:Num_donations,data=t.
  family=binomial)
glm.probs=predict(glm.fit,test_data,type="response")
out=data.frame(test_data$ID,glm.probs)
names(out)=c("", "Made Donation in March 2007")
write.csv(out,file="logreg_int_2018_04_16.csv",row.names=FALSE)

# fit KNN, choose k to minimize the cost on the validation set
library(class)
train.X=as.matrix(train_data[, -c(1,4,6)],nrow=nrow(train_data),ncol=ncol(train_data)-3)
val.X=as.matrix(val_data[, -c(1,4,6)],nrow=nrow(val_data),ncol=ncol(val_data)-3)
set.seed(111)
kvals=seq(1,nrow(train_data))
logloss=rep(NA,length(kvals))
for (k in 1:length(kvals)){
  knn.pred=knn(train.X,val.X,train_data$Made_donation,k=k,prob=TRUE)
  knn.prob=abs(attr(knn.pred,"prob")-10^-10) # need to add a fudge factor to deal
  # with logs of 0 and 1
  logloss[k]=-(1/nrow(val_data))*( sum(val_data$Made_donation*log(knn.prob)) +
    sum((1-val_data$Made_donation)*log(1-knn.prob)) )
}
which.min(logloss)


## [1] 316

logloss[which.min(logloss)]

## [1] 1.120105

plot(1:length(kvals),logloss)

```

```
# choose k at the elbow, note small k is high variance, low bias:
```

```
plot(1:50,logloss[1:50])
```

```
# calculate predictions using all of the training data
```

```
k=27 # pick k = 27
```

```
final_train.X=rbind(train.X,val.X)
```

```
test.X=as.matrix(test_data[,-c(1,4)],nrow=nrow(test_data),ncol=ncol(test_data)-2)
```

```
knn.pred=knn(final_train.X,test.X,c(train_data$Made_donation,val_data$Made_donation),k=k,  
             prob=TRUE)
```

```
knn.prob=abs(attr(knn.pred,"prob")-10^-10)
```

```
out=data.frame(test_data$ID,knn.prob)
```

```
names(out)=c("", "Made Donation in March 2007")
```

```
write.csv(out,file="knn_2018_04_16.csv",row.names=FALSE)
```

```
# fit a random forest with bagging
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.4.4
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
set.seed(200)
```

```
train_data$Made_donation=as.factor(train_data$Made_donation)
```

```
val_data$Made_donation=as.factor(val_data$Made_donation)
```

```
#test_data$Made_donation=as.factor(test_data$Made_donation)
```

```
m=ncol(train_data)-3
```

```
bag.data=randomForest(Made_donation~Last_donation+Num_donations+First_donation+Last_donation:Num_donations)
```

```
bag.data
```

```
##
## Call:
## randomForest(formula = Made_donation ~ Last_donation + Num_donations +      First_donation + Last_d
##               Type of random forest: classification
##               Number of trees: 500
## No. of variables tried at each split: 3
##
## OOB estimate of error rate: 24.81%
## Confusion matrix:
##      0  1 class.error
## 0 271 37   0.1201299
## 1  63 32   0.6631579

importance(bag.data)


##              0              1 MeanDecreaseAccuracy MeanDecreaseGini
## Last_donation  7.907709 11.659565             13.06505             31.59352
## Num_donations 21.943140  9.541482             26.45239             41.02613
## First_donation 18.246722 -1.557938             17.11908             51.15227

probs=predict(bag.data,newdata=val_data,type="prob")[,2]+10^-10
# log loss
-(1/nrow(val_data))*( sum((as.numeric(val_data$Made_donation)-1)*log(probs)) +
  sum((1-(as.numeric(val_data$Made_donation)-1))*log(1-probs)) )

## [1] 1.13197

# fit a random forest with boosting
library(gbm)

## Warning: package 'gbm' was built under R version 3.4.4
## Loading required package: survival
## Loading required package: lattice
## Loading required package: splines
## Loading required package: parallel
## Loaded gbm 2.1.3
```



```

lambdas=c(10^-5,10^-4,10^-3,10^-2,0.05,0.1,0.2,0.5)
logloss=rep(NA,length(lambdas))
train_data_exp=cbind(train_data,train_data$Last_donation*train_data$Num_donations)
names(train_data_exp)[7]="Last_num_int"
val_data_exp=cbind(val_data,val_data$Last_donation*val_data$Num_donations)
names(val_data_exp)[7]="Last_num_int"
for (i in 1:length(lambdas)){
  lambda=lambdas[i]
  boost.data=gbm(Made_donation~Last_donation+Num_donations+First_donation+Last_num_int,data=train_data,
                 distribution="bernoulli",n.trees=1000,shrinkage=lambda,interaction.depth=2)
  probs=predict(bag.data,newdata=val_data_exp,type="prob")[,2]+10^-10
  logloss[i]=-(1/nrow(val_data))* ( sum((as.numeric(val_data$Made_donation)-1)*log(probs)) +
    sum((1-(as.numeric(val_data$Made_donation)-1))*log(1-probs)) )
}
logloss

## [1] 1.13197 1.13197 1.13197 1.13197 1.13197 1.13197 1.13197 1.13197

```