

Blood Donations

Bridget Bertoni

April 12, 2018

```
# read in data
train_data=read.csv(file="/Users/bbertoni/Desktop/github/Blood_Donations/training_data.csv",header=T)
test_data=read.csv(file="/Users/bbertoni/Desktop/github/Blood_Donations/test_data.csv",header=T)

names(train_data)=c("ID","Last_donation","Num_donations","Volume","First_donation","Made_donation")
names(test_data)=c("ID","Last_donation","Num_donations","Volume","First_donation")

head(train_data)

##      ID Last_donation Num_donations Volume First_donation Made_donation
## 1 619              2             50  12500             98             1
## 2 664              0             13   3250             28             1
## 3 441              1             16   4000             35             1
## 4 160              2             20   5000             45             1
## 5 358              1             24   6000             77             0
## 6 335              4              4   1000              4             0

head(test_data)

##      ID Last_donation Num_donations Volume First_donation
## 1 659              2             12   3000             52
## 2 276              21              7   1750             38
## 3 263              4              1    250              4
## 4 303              11             11   2750             38
## 5  83              4             12   3000             34
## 6 500              3             21   5250             42

# split training data into a training set and a validation set
set.seed(333)
train=sample(1:nrow(train_data),0.7*nrow(train_data),replace=F)
val_data = train_data[-train,]
train_data = train_data[train,]

# check for missing or strange values
sum(is.na(train_data))

## [1] 0

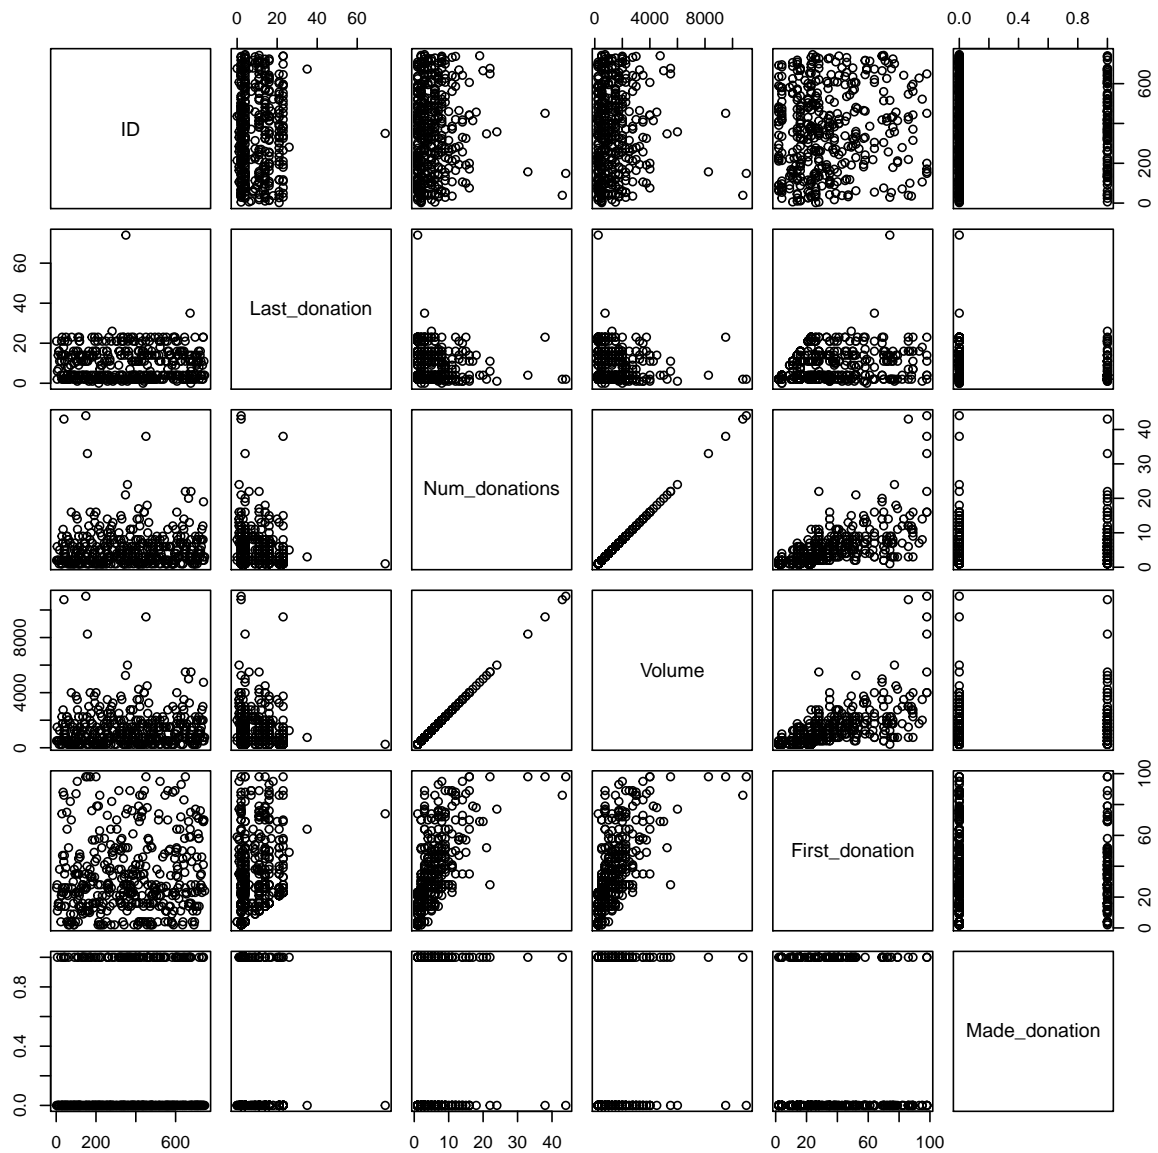
sum(is.na(val_data))# no missing values

## [1] 0

sum(is.na(test_data)) # missing values

## [1] 0

plot(train_data)
```



```
cor(train_data) # correlation between volume and the number of donations is 1
```

```
##           ID Last_donation Num_donations      Volume
## ID      1.00000000    0.02180929    0.02167073    0.02167073
## Last_donation 0.02180929    1.00000000   -0.16150193 -0.16150193
## Num_donations 0.02167073   -0.16150193    1.00000000    1.00000000
## Volume      0.02167073   -0.16150193    1.00000000    1.00000000
## First_donation 0.10391260    0.17791940    0.64740429    0.64740429
## Made_donation 0.04101195   -0.24721682    0.20707525    0.20707525
##           First_donation Made_donation
## ID      0.103912600    0.041011949
## Last_donation 0.177919401 -0.247216821
## Num_donations 0.647404293    0.207075245
## Volume     0.647404293    0.207075245
```

```
## First_donation      1.000000000    0.008180282
## Made_donation       0.008180282    1.000000000

# fit basic logistic regression
glm.fit=glm(Made_donation~Last_donation+Num_donations+First_donation,data=train_data,
            family=binomial)
summary(glm.fit)

##
## Call:
## glm(formula = Made_donation ~ Last_donation + Num_donations +
##      First_donation, family = binomial, data = train_data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3459  -0.7959  -0.5454  -0.3334   2.4440
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -0.715501   0.240869  -2.970  0.00297 **
## Last_donation -0.082632   0.021242  -3.890  0.00010 ***
## Num_donations  0.109057   0.033638   3.242  0.00119 **
## First_donation -0.012572   0.007533  -1.669  0.09513 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 440.16  on 402  degrees of freedom
## Residual deviance: 396.46  on 399  degrees of freedom
## AIC: 404.46
##
## Number of Fisher Scoring iterations: 5

#plot(glm.fit)
glm.probs=predict(glm.fit,val_data,type="response")
glm.pred=rep(0,nrow(val_data))
glm.pred[glm.probs>0.5]=1
table(glm.pred,val_data$Made_donation) # confusion matrix

##
## glm.pred    0    1
##           0 129  38
##           1   1   5

mean(glm.pred==val_data$Made_donation)

## [1] 0.7745665

# fit KNN, choose k to minimize the cost on the validation set
library(class)
train.X=as.matrix(train_data,nrow=nrow(train_data),ncol=ncol(train_data))
val.X=as.matrix(val_data,nrow=nrow(val_data),ncol=ncol(val_data))
set.seed(111)
knn.pred=knn(train.X,val.X,train_data$Made_donation,k=1)

table(knn.pred,val_data$Made_donation)
```

```
##  
## knn.pred    0    1  
##           0 100  30  
##           1  30  13  
mean(knn.pred==val_data$Made_donation)  
  
## [1] 0.6531792  
# fit a random forest
```